

Kubebuilder

Tutorial

Ondrej Sery

osery@purestorage.com



PureStorage

Delivering the Modern Data Experience

GROWTH



+\$50B

Total Addressable Market

\$1.64B

FY20 Annual Revenue

Subscription Services up **37%** Year-over-year (Q1 FY21)

CUSTOMERS



7,800+

Customers

1,700

New Customers in FY20


45%

of Fortune 500
Companies

33%+


of Bookings from Cloud
Customers (Q1 FY21)

LEADERSHIP



Gartner Magic Quadrant

For the sixth year in a row,
Pure Storage is a leader in the
Gartner Magic Quadrant.



Net Promoter Score

In the top 1% of B2B companies.

COMPANY MILESTONES



2009

Founded

2011

Pioneer in
All-Flash

2012

Introduction of
FlashArray™

Late 2012

First
international
office open in
Europe

Mid 2014

Evergreen™
Storage
Non-Disruptive
Upgrades

Early 2015

Pure1®
Cloud-Based
Management

Late 2015

IPO (PSTG)

2016

Introduction of
FlashBlade®

Early 2018

Pure introduces
the first AI-ready
Infrastructure
(AIRI™)

Mid 2018

Launch of as-a-
Service model
(Pure as-a-Service)

2018

Cloud Data
Services

2018

\$1+ Billion in
revenue

2019

Compuverde
Acquisition

2019

Expands
FlashArray™
Family

2019

Named a
Leader:
Gartner Magic
Quadrant

2020

Pure as-a-Service
celebrates 2-year
anniversary



Kubernetes Operators

THE WHY

“Operating” Kubernetes Applications

- **Stateless applications**

- Easy enough using Kubernetes alone
- New version & rolling update

- **Stateful applications**

- R/O mode, Backup, Migration, Upgrade, R/W mode
- Manually by a human “operator” with domain knowledge

⇒ *Automated “operator” with domain knowledge*

Custom Resource Definitions

- New Kubernetes resources/Kinds
 - Domain specific description of the desired state.
 - e.g., Prometheus, PrometheusRule

```
apiVersion: monitoring.coreos.com/v1
kind: Prometheus
metadata:
  name: prometheus
spec:
  serviceAccountName: prometheus
  serviceMonitorSelector:
    matchLabels:
      team: frontend
  resources:
    requests:
      memory: 400Mi
  enableAdminAPI: true
```

```
apiVersion: monitoring.coreos.com/v1
kind: PrometheusRule
metadata:
  labels:
    prometheus: example
    role: alert-rules
  name: prometheus-example-rules
spec:
  groups:
    - name: ./example.rules
      rules:
        - alert: ExampleAlert
          expr: vector(1)
```

Kubernetes Control Loop

OBSERVE → CHECK DIFFERENCE → TAKE ACTION

- Observe:
 - Kubernetes resources/Kinds with the desired state
 - *E.g., replica set with 3 replicas*
- Check Difference:
 - Between desired and actual state
 - *E.g., only 2 running pods*
- Take Action:
 - To remove the difference
 - *E.g., start another pod*

Tooling

- **Kubebuilder**
- Operator SDK
 - Helm
 - Ansible
 - Golang (now using Kubebuilder)
- KUDO
- Metacontroller

Kubebuilder

THE HOW

Kubebuilder Overview

- Generates a project skeleton for an operator
 - Go schema & API definitions
 - Go controller skeletons
 - Deployment yaml files
 - Test skeletons
 - Dockerfile
 - Makefile
- Covers also evolution
 - Add new (versions of) resources

Tutorial: Coffee Maker Operator

- **Existing:** Coffee Maker App
 - GitHub: github.com/osery/coffee-maker
 - Docker image: ghcr.io/osery/coffee-maker:latest
- **Our goal:** *Coffee Maker K8s Operator*
 - CRD-based API for the Coffee Maker App
 - Controller
 - Call the coffee maker REST API
 - New beverage
 - Status updates

New Kubebuilder Project

- We need a new go module:
`$ go mod init coffee.demo.purestorage.com`
- Generate the project skeleton:
`$ kubebuilder init --domain coffee.demo.purestorage.com`

Let's try it and explore a bit...

New API

- Let's create a Coffee CRD:

```
$ kubebuilder create api --group beverage --version v1 --kind Coffee
```

- *Generates:*
 - Golang custom resource types with json bindings
 - K8s CRD and RBAC yaml definitions
 - Example custom resource yaml files
 - Controller skeleton
 - Boilerplate registration code, etc.

Add Fields

- Manual edits to add/remove/update the actual payload
- Special types:
 - `resource.Quantity` (e.g., “1.5Gi”)
 - `metav1.Time` (`time.Time` with correct yaml/json marshalling)
- Our “Coffee” fields:
 - **Spec:**
 - `type = {espresso, americano, latte}`
 - `extraSugar = {true, false}`
 - **Status:**
 - `status = {queued, brewing, done, failed}`

Build and Deploy CRDs

- Regenerate everything and build:
`$ make manifests`
- Install CRD definitions into a K8s cluster (current kubectl default)
`$ make install`
- Uninstall CRD definitions
`$ make uninstall`

Customizing CRD

- Mark additional fields in the default "kubectl get" display:

```
// +kubebuilder:printcolumn:name="Type",type=string,JSONPath=`.spec.type`  
// +kubebuilder:printcolumn:name="ExtraSugar",type=boolean,JSONPath=`.spec.extraSugar`  
// +kubebuilder:printcolumn:name="Status",type=string,JSONPath=`.status.status`
```

- Ensure only allowed values are used as coffee Type.

```
// +kubebuilder:validation:Enum=espresso;latte;americano
```


Controller

- Get K8s resource
- Get REST resource
- Create a new if missing
- Update Status
- Requeue until in status **done**

Build and Deploy Controller

- Build a push docker image:
`$ make docker-build docker-push IMG=repository/image:tag`
- Install the controller into a K8s cluster
`$ make deploy IMG=repository/image:tag`

RESOURCES

- Kubebuilder
 - GitHub: <https://github.com/kubernetes-sigs/kubebuilder>
 - Book: <https://kubebuilder.io>
- K3s: <https://k3s.io>
- Demo materials:
 - Will upload to <https://github.com/osery/2020-11-24-kubebuilder-tutorial> after the talk...

Q&A